Oasis Infobyte Data Science Internship

Task 3 - Sales Prediction Using Python

Ranjeeta Kumari

In [1]:

```python
#importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

In [7]:

```python
#loading the dataset
ad=pd.read_csv(r'C:\Users\ajeet singh\Downloads\Advertising.csv')
print(ad)
```

```
     Unnamed: 0      TV  Radio  Newspaper  Sales
0             1   230.1   37.8       69.2   22.1
1             2    44.5   39.3       45.1   10.4
2             3    17.2   45.9       69.3    9.3
3             4   151.5   41.3       58.5   18.5
4             5   180.8   10.8       58.4   12.9
..          ...     ...    ...        ...    ...
195         196    38.2    3.7       13.8    7.6
196         197    94.2    4.9        8.1    9.7
197         198   177.0    9.3        6.4   12.8
198         199   283.6   42.0       66.2   25.5
199         200   232.1    8.6        8.7   13.4

[200 rows x 5 columns]
```

In [8]:

```python
ad.head()
```

Out[8]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

```
ad.tail()
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

```
ad.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  ------      --------------  -----
 0   Unnamed: 0  200 non-null    int64
 1   TV          200 non-null    float64
 2   Radio       200 non-null    float64
 3   Newspaper   200 non-null    float64
 4   Sales       200 non-null    float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
```

```
ad.describe()
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000	200.000000
mean	100.500000	147.042500	23.264000	30.554000	14.022500
std	57.879185	85.854236	14.846809	21.778621	5.217457
min	1.000000	0.700000	0.000000	0.300000	1.600000
25%	50.750000	74.375000	9.975000	12.750000	10.375000
50%	100.500000	149.750000	22.900000	25.750000	12.900000
75%	150.250000	218.825000	36.525000	45.100000	17.400000
max	200.000000	296.400000	49.600000	114.000000	27.000000

In [16]:

```python
#scaling the features
scaler = StandardScaler()
ad[['TV','Radio','Newspaper']] = scaler.fit_transform(ad[['TV','Radio','Newspaper']])
```

In [18]:

```python
X = ad.drop(['Sales'], axis=1)
y = ad['Sales']

x_train, x_test, y_train, y_test = train_test_split(X,y, test_size=0.2, random_state=100
```

In [19]:

```python
from sklearn.linear_model import LinearRegression
```

In [20]:

```python
#LINEAR REGRESSION ALGORITHM
lr = LinearRegression()
#fitting the model
lr.fit(x_train,y_train)
y_pred = lr.predict(x_test)
```

In [21]:

```python
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

In [22]:

```python
print(mean_absolute_error(y_true=y_test, y_pred=y_pred))
print(mean_squared_error(y_true=y_test, y_pred=y_pred))
```

```
1.029124930516935
1.733787988439963
```

In [23]:

```python
from sklearn.metrics import r2_score
```

In [24]:

```python
#printing the accuracy of the model
d=r2_score(y_true=y_test, y_pred=y_pred)*100
print(f'THE ACCURACY SCORE IS : {d} %')
```

```
THE ACCURACY SCORE IS : 91.84136004105717 %
```